

# CS373 Data Mining and Machine Learning

## Lecture 8

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

*(originally prepared by Tommi Jaakkola, MIT CSAIL)*

# Today's topics

- Ensembles and Boosting
  - ensembles, relation to feature selection
  - myopic forward-fitting and boosting
  - understanding boosting

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Ensembles

- An ensemble classifier combines a set of  $m$  “weak” base learners into a “strong” ensemble

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

ensemble  
discriminant  
function

non-negative  
“votes”

base learner

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Ensembles

- An ensemble classifier combines a set of  $m$  “weak” base learners into a “strong” ensemble

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$


ensemble  
discriminant  
function

non-negative  
“votes”

Assignment Project Exam Help  
<https://powcoder.com>

base learner

- The base learners are typically simple “decision stumps”, i.e., linear classifiers based on one coordinate



# Ensembles

- An ensemble classifier combines a set of  $m$  “weak” base learners into a “strong” ensemble

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$


ensemble  
discriminant  
function

non-negative  
“votes”

Assignment Project Exam Help  
<https://powcoder.com>

base learner

- The base learners are typically simple “decision stumps”, i.e., linear classifiers based on one coordinate



# Ensembles

- An ensemble classifier combines a set of  $m$  “weak” base learners into a “strong” ensemble

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$


ensemble  
discriminant  
function

non-negative  
“votes”

Assignment Project Exam Help  
<https://powcoder.com>

base learner

- The base learners are typically simple “decision stumps”, i.e., linear classifiers based on one coordinate



# Ensemble learning

- We can view the ensemble learning problem as a coordinate selection problem

$$h_m(\underline{x}) = \begin{bmatrix} 0 \\ \vdots \\ \alpha_m \\ 0 \\ \vdots \end{bmatrix} \begin{bmatrix} 0 \\ \vdots \\ h(\underline{x}; \underline{\theta}_1) \\ \vdots \\ h(\underline{x}; \underline{\theta}_m) \\ \vdots \end{bmatrix}$$

parameters  
or “votes”
coordinates  
indexed by  $\underline{\theta}$

- The problem of finding the best “coordinates” to include corresponds to finding the parameters of the base learners (out of an uncountable set)

# Estimation criterion

- In principle, we can estimate the ensemble

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$


by minimizing the training loss

$$\sum_{t=1}^n \text{Loss}(y_t h_m(\underline{x}_t))$$


Assignment Project Exam Help  
<https://powcoder.com>

with respect to the parameters in the ensemble

logistic loss



exponential loss





# Estimation criterion

- In principle, we can estimate the ensemble

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

by minimizing the training loss

$$\sum_{t=1}^n \text{Loss}(y_t, h_m(\underline{x}_t))$$

with respect to the parameters in the ensemble

- This is a hard problem to solve jointly but we can add base learners sequentially (cf. forward fitting)

# Estimation criterion

- In principle, we can estimate the ensemble

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

by minimizing the training loss

$$\sum_{t=1}^n \text{Loss}(y_t h_m(\underline{x}_t))$$

Assignment Project Exam Help  
<https://powcoder.com>

with respect to the parameters in the ensemble

Add WeChat powcoder

- This is a hard problem to solve jointly but we can add base learners sequentially (cf. forward fitting)

Fix  $h_{m-1}(\underline{x})$

Find  $\alpha_m$  and  $\underline{\theta}_m$  that minimize  $y_t h_m(\underline{x}_t)$

$$J(\alpha_m, \underline{\theta}_m) = \sum_{t=1}^n \text{Loss} \left( \underbrace{y_t h_{m-1}(\underline{x}_t)}_{\text{fixed}} + \alpha_m y_t h(\underline{x}_t; \underline{\theta}_m) \right)$$

# Myopic forward-fitting

$$J(\alpha_m, \underline{\theta}_m) = \sum_{t=1}^n \text{Loss} \left( \overbrace{y_t h_{m-1}(\underline{x}_t) + \alpha_m y_t h(\underline{x}_t; \underline{\theta}_m)}^{y_t h_m(\underline{x}_t)} \right)$$

fixed

- Out of all  $\underline{\theta}_m$  we wish to select one that minimizes the derivative of the loss at  $\alpha_m = 0$  (has the most negative derivative at zero)

$$\begin{aligned} \left. \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m=0} &= \sum_{t=1}^n \left[ \left. \frac{\partial}{\partial z} \text{Loss}(z) \right|_{z=y_t h_{m-1}(\underline{x}_t)} \right] y_t h(\underline{x}_t; \underline{\theta}_m) \\ &= \sum_{t=1}^n \text{DLoss} \left( \underbrace{y_t h_{m-1}(\underline{x}_t)}_{\text{fixed weights on training examples}} \right) y_t h(\underline{x}_t; \underline{\theta}_m) \end{aligned}$$

# Weights on training examples

$$\begin{aligned}\frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \Big|_{\alpha_m=0} &= \sum_{t=1}^n \left[ \frac{\partial}{\partial z} \text{Loss}(z) \Big|_{z=y_t h_{m-1}(\underline{x}_t)} \right] y_t h(\underline{x}_t; \underline{\theta}_m) \\ &= \sum_{t=1}^n \text{DLoss}(y_t h_{m-1}(\underline{x}_t)) y_t h(\underline{x}_t; \underline{\theta}_m)\end{aligned}$$

Assignment Project Exam Help  
these derivatives are negative

<https://powcoder.com>

Add WeChat powcoder

# Weights on training examples

$$\left. \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m=0} = \sum_{t=1}^n \left[ \frac{\partial}{\partial z} \text{Loss}(z) \Big|_{z=y_t h_{m-1}(\underline{x}_t)} \right] y_t h(\underline{x}_t; \underline{\theta}_m)$$

$$= \sum_{t=1}^n \text{DLoss}(y_t h_{m-1}(\underline{x}_t)) y_t h(\underline{x}_t; \underline{\theta}_m)$$

Assignment Project Exam Help  
these derivatives are negative

$$= \sum_{t=1}^n W_t (-y_t) h(\underline{x}_t; \underline{\theta}_m)$$

https://powcoder.com  
Add WeChat powcoder  
positive weight      error (agreement with the opposite label)

# Weights on training examples

$$\left. \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m=0} = \sum_{t=1}^n \left[ \left. \frac{\partial}{\partial z} \text{Loss}(z) \right|_{z=y_t h_{m-1}(\underline{x}_t)} \right] y_t h(\underline{x}_t; \underline{\theta}_m)$$

$$= \sum_{t=1}^n \text{DLoss}(y_t h_{m-1}(\underline{x}_t)) y_t h(\underline{x}_t; \underline{\theta}_m)$$

Assignment Project Exam Help

$$= \sum_{t=1}^n W_t (-y_t) h(\underline{x}_t; \underline{\theta}_m)$$

these derivatives are negative  
<https://powcoder.com>  
 Add WeChat powcoder  
 error (agreement with  
 positive weight the opposite label)

Logistic loss:  $W_t = g(-y_t h_{m-1}(\underline{x}_t)), \quad g(z) = (1 + \exp(-z))^{-1}$

Exponential loss:  $W_t = \exp(-y_t h_{m-1}(\underline{x}_t))$

# Weights on training examples

$$\begin{aligned}
 \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \Big|_{\alpha_m=0} &= \sum_{t=1}^n \left[ \frac{\partial}{\partial z} \text{Loss}(z) \Big|_{z=y_t h_{m-1}(\underline{x}_t)} \right] y_t h(\underline{x}_t; \underline{\theta}_m) \\
 &= \sum_{t=1}^n \underbrace{\text{DLoss}(y_t h_{m-1}(\underline{x}_t))}_{\text{these derivatives are negative}} y_t h(\underline{x}_t; \underline{\theta}_m) \\
 &= \sum_{t=1}^n \underbrace{W_t (-y_t)}_{\substack{\text{Add WeChat powcoder} \\ \text{error (agreement with} \\ \text{positive weight} \quad \text{the opposite label)}}} h(\underline{x}_t; \underline{\theta}_m)
 \end{aligned}$$

Logistic loss:  $W_t = g(-y_t h_{m-1}(\underline{x}_t)), \quad g(z) = (1 + \exp(-z))^{-1}$

Exponential loss:  $W_t = \exp(-y_t h_{m-1}(\underline{x}_t))$

- We can always normalize the weights without affecting the choice of  $\underline{\theta}_m$

$$W_t \leftarrow \frac{W_t}{\sum_{i=1}^n W_i}$$

# General boosting algorithm

- We use a myopic forward-fitting method to estimate

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

Step 0:  $h_0(\underline{x}) = 0$ ,  $W_t = 1/n$ ,  $t = 1, \dots, n$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# General boosting algorithm

- We use a myopic forward-fitting method to estimate

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

Step 0:  $h_0(\underline{x}) = 0$ ,  $W_t = 1/n$ ,  $t = 1, \dots, n$

Step 1: Find  $\hat{\underline{\theta}}_m$  that minimizes the weighted error

$$\left. \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m=0} = \sum_{t=1}^n W_t (-y_t) h(\underline{x}_t; \underline{\theta}_m)$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# General boosting algorithm

- We use a myopic forward-fitting method to estimate

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

Step 0:  $h_0(\underline{x}) = 0$ ,  $W_t = 1/n$ ,  $t = 1, \dots, n$

Step 1: Find  $\hat{\underline{\theta}}_m$  that minimizes the weighted error

$$\left. \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m=0} = \sum_{t=1}^n W_t (-y_t) h(\underline{x}_t; \underline{\theta}_m)$$

- weights correspond to derivatives of the loss function
- the weight is large if the example is not classified correctly by the ensemble we have so far
- finding the parameters that minimize the weighted error is easily solved, e.g., for decision stumps

# General boosting algorithm

- We use a myopic forward-fitting method to estimate

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

Step 0:  $h_0(\underline{x}) = 0$ ,  $W_t = 1/n$ ,  $t = 1, \dots, n$

Step 1: Find  $\hat{\underline{\theta}}_m$  that minimizes the weighted error

$$\left. \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m=0} = \sum_{t=1}^n W_t (-y_t) h(\underline{x}_t; \underline{\theta}_m)$$

Step 2: Find  $\hat{\alpha}_m$  that minimizes

$$J(\alpha_m, \hat{\underline{\theta}}_m) = \sum_{t=1}^n \text{Loss} \left( \underbrace{y_t h_{m-1}(\underline{x}_t)}_{\text{fixed}} + \alpha_m \underbrace{y_t h(\underline{x}_t; \hat{\underline{\theta}}_m)}_{\text{fixed}} \right)$$

- this is a 1-dimensional convex problem that can be solved easily.

# General boosting algorithm

- We use a myopic forward-fitting method to estimate

$$h_m(\underline{x}) = \alpha_1 h(\underline{x}; \underline{\theta}_1) + \dots + \alpha_m h(\underline{x}; \underline{\theta}_m)$$

Step 0:  $h_0(\underline{x}) = 0$ ,  $W_t = 1/n$ ,  $t = 1, \dots, n$

Step 1: Find  $\hat{\underline{\theta}}_m$  that minimizes the weighted error

$$\left. \frac{\partial J(\alpha_m, \underline{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m=0} = \sum_{t=1}^n W_t (-y_t) h(\underline{x}_t; \underline{\theta}_m)$$

Step 2: Find  $\hat{\alpha}_m$  that minimizes

$$J(\alpha_m, \hat{\underline{\theta}}_m) = \sum_{t=1}^n \text{Loss} \left( \underbrace{y_t h_{m-1}(\underline{x}_t)}_{\text{fixed}} + \alpha_m \underbrace{y_t h(\underline{x}_t; \hat{\underline{\theta}}_m)}_{\text{fixed}} \right)$$

Step 3: Update example weights

$$W_t = -\text{DLoss} \left( \underbrace{y_t h_{m-1}(\underline{x}_t) + \hat{\alpha}_m y_t h(\underline{x}_t; \hat{\underline{\theta}}_m)}_{y_t h_m(\underline{x}_t)} \right)$$


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

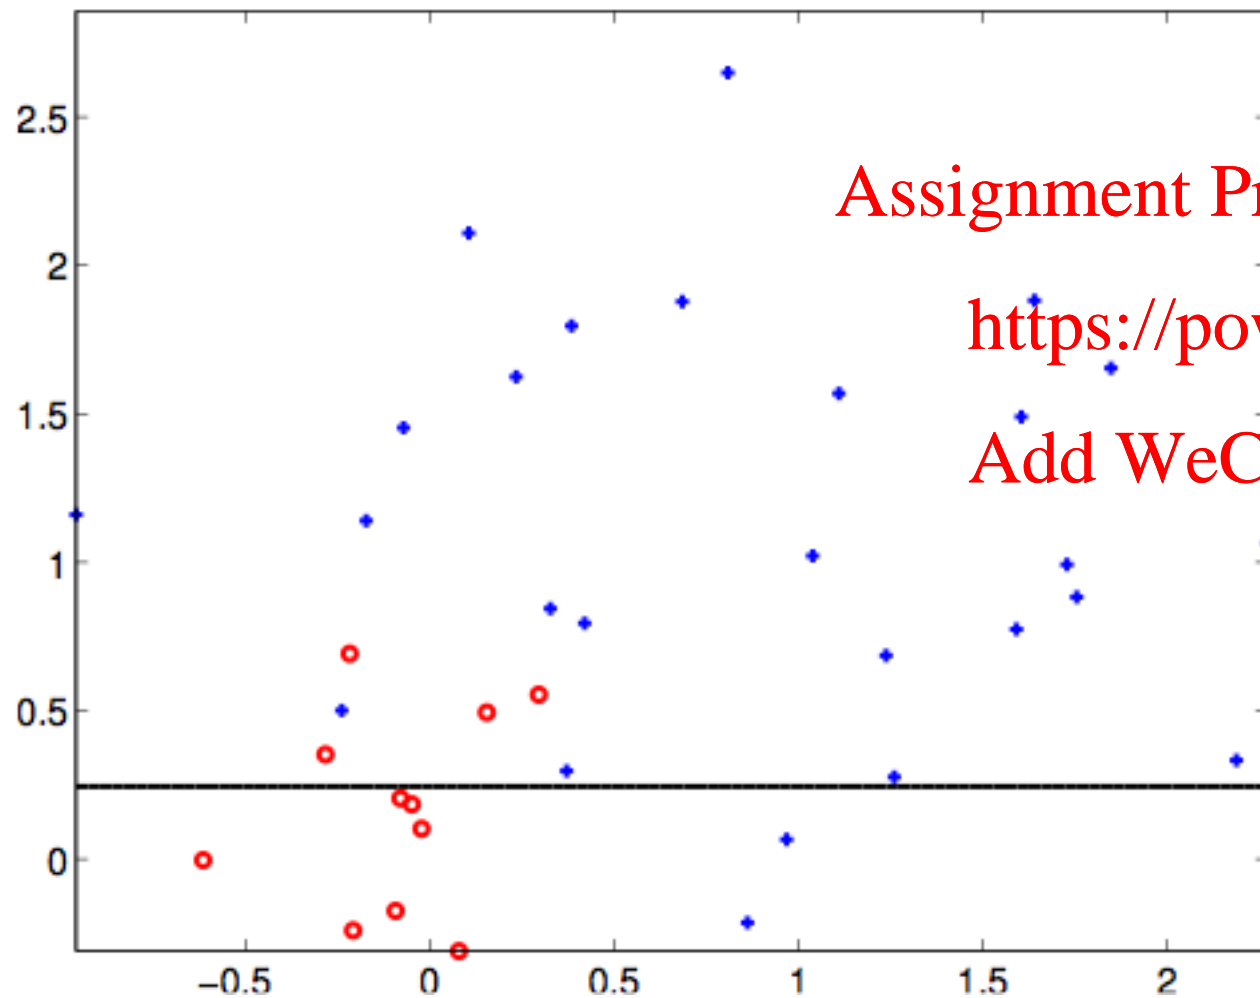
# Boosting example

Logistic loss  $\text{Loss}(z) = \log(1 + \exp(-z))$




# Boosting example

$$h(\underline{x}; \hat{\theta}_1)$$



$$h_1(\underline{x}) = \hat{\alpha}_1 h(\underline{x}; \hat{\theta}_1)$$




Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

# Boosting example

$$h(\underline{x}; \hat{\theta}_2)$$



$$h_2(\underline{x}) = \hat{\alpha}_1 h(\underline{x}; \hat{\theta}_1) + \hat{\alpha}_2 h(\underline{x}; \hat{\theta}_2)$$




Assignment Project Exam Help

<https://powcoder.com>

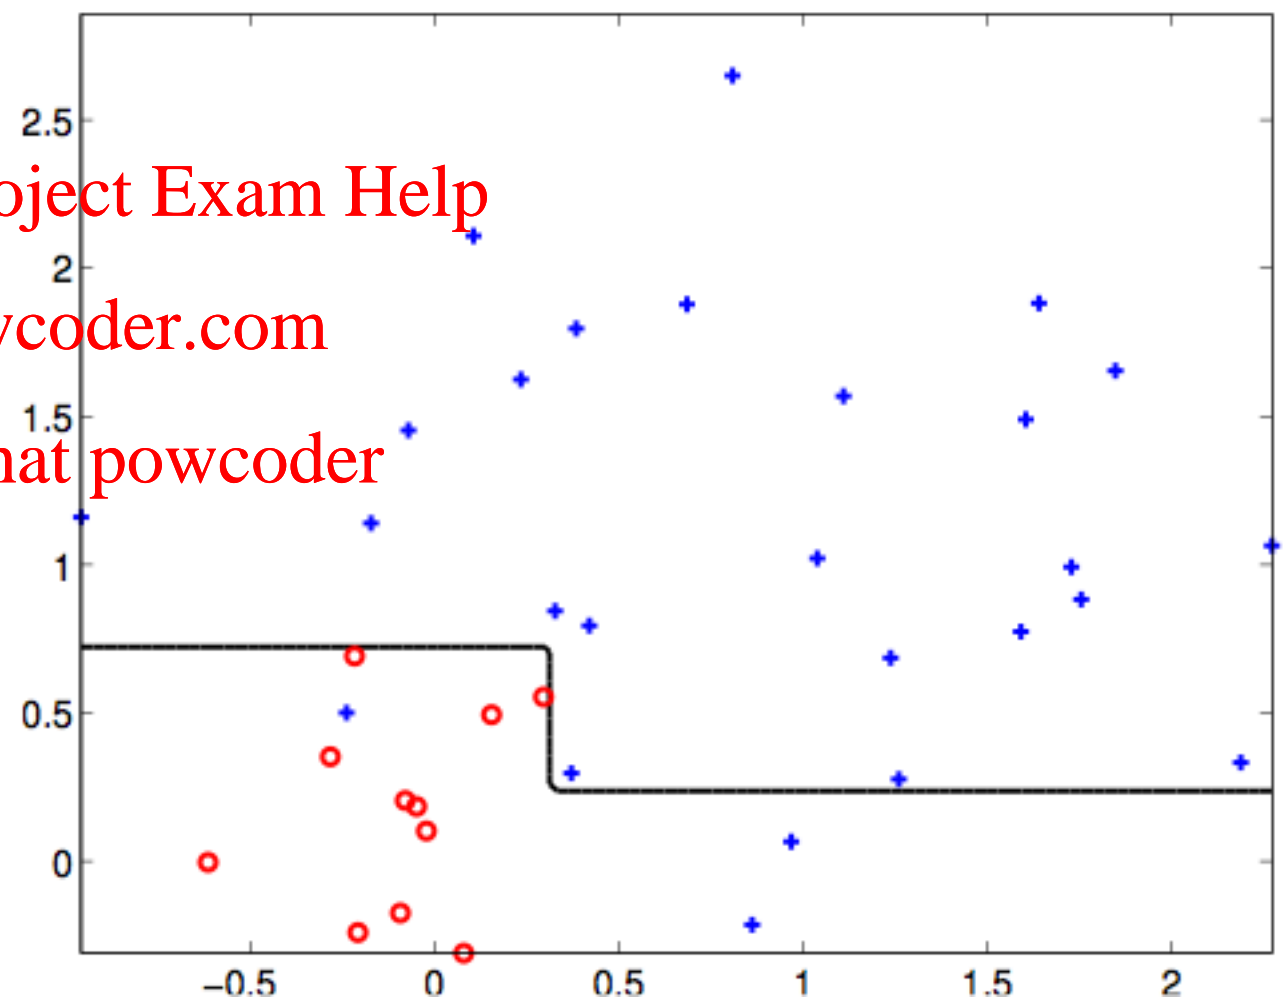
Add WeChat powcoder

# Boosting example

$$h(\underline{x}; \hat{\theta}_3)$$



$$h_3(\underline{x}) = \hat{\alpha}_1 h(\underline{x}; \hat{\theta}_1) + \cdots + \hat{\alpha}_3 h(\underline{x}; \hat{\theta}_3)$$



Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder




# Boosting example

$$h(\underline{x}; \hat{\theta}_4)$$



$$h_4(\underline{x}) = \hat{\alpha}_1 h(\underline{x}; \hat{\theta}_1) + \cdots + \hat{\alpha}_4 h(\underline{x}; \hat{\theta}_4)$$



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Trying the same base learner again

- $\hat{\alpha}_m$  is set to minimize the loss given the base learner

$$J(\alpha_m, \hat{\theta}_m) = \sum_{t=1}^n \text{Loss}(y_t h_{m-1}(\underline{x}_t) + \alpha_m y_t h(\underline{x}_t; \hat{\theta}_m))$$

- At the optimum value,

Assignment Project Exam Help

$$\left. \frac{\partial J(\alpha_m, \hat{\theta}_m)}{\partial \alpha_m} \right|_{\alpha_m = \hat{\alpha}_m}$$

<https://powcoder.com>

Add WeChat powcoder

$$\sum_{t=1}^n \text{DLoss}(y_t h_{m-1}(\underline{x}_t) + \hat{\alpha}_m y_t h(\underline{x}_t; \hat{\theta}_m)) y_t h(\underline{x}_t; \hat{\theta}_m) = 0$$

updated weights (up to  
normalization and overall sign)

- Thus the current base learner is useless at the next iteration (but may be chosen again later on)

# Ensemble training error

- The boosting algorithm decreases the training loss

$$\sum_{t=1}^n \text{Loss}(y_t h_m(\underline{x}_t))$$

monotonically while the base learners remain effective against the weighted error (derivative is not zero)

- For any non-increasing loss function,

$$\sum_{t=1}^n I_{[y_t h_m(\underline{x}_t) \leq 0]} \leq \frac{1}{\text{Loss}(0)} \sum_{t=1}^n \text{Loss}(y_t h_m(\underline{x}_t))$$

Thus we have a monotonically decreasing upper bound on the 0-1 training error (classification error)

# Ensemble training error

